

# Understand `std::vector`

Huỳnh Ngọc Thi  
so61pi.re@gmail.com

# Why?

- `std::vector` is a fundamental data structure in C++.
- It's used everywhere.
- It's array-based, good for performance.
  - Maybe not the best, but still.
- Its iterator's category is `RandomAccessIterator`, which is required by some algorithms
  - `shuffle`, `sort`, `stable_sort`, `nth_element`, `make_heap`, `sort_heap`, ...

# Placement new

- How can we construct `std::string` on a given buffer?

- Placement new

```
T* = new (buffer) T;  
char buffer[sizeof(string)] = {};  
auto pstr = new (buffer) string{"abc"};
```

- How can we destruct it?

- Placement delete?
  - We must call destructor ourselves.

```
pstr->~string();
```

- <https://isocpp.org/wiki/faq/dtors#placement-new>


# std::vector memory layout



- Why do we have extra “unused” space?
  - `push_back/emplace_back` can have amortized constant complexity.
- Amortized constant complexity
  - “The basic idea is that a worst case operation can alter the state in such a way that the worst case cannot occur again for a long time, thus “amortizing” its cost.”
  - [https://en.wikipedia.org/wiki/Amortized\\_analysis](https://en.wikipedia.org/wiki/Amortized_analysis)

# std::vector memory layout

- Basically, std::vector is constructed from 3 pointers.



# push\_back complexity analysis

- If we don't have that “unused” space
  - Each push\_back
    - allocate new memory
    - move/copy old elements
    - construct new element
    - destruct old elements
    - free old memory.
  - $O(n)$

# push\_back complexity analysis

- Else,
  - New element can be constructed without reallocating memory.
- We have a growth factor,  $k$ .
  - Each time reallocation happens, new size =  $\max(\text{size}() + 1, \text{size}() * k)$ .
- MSVC uses 1.5 (`_Grow_to`), libstdc++ 2 (`_M_check_len`).


# push\_back complexity analysis

k = 2				
before push_back		after push_back		reallocation
size()	capacity()	size()	capacity()	
0	0	1	1	true
1	1	2	2	true
2	2	3	4	true
3	4	4	4	false
...	...	...	...	...
127	128	128	128	false
128	128	129	256	true
129	256	130	256	false
...	...	...	...	...



# push\_back complexity analysis

- If we push\_back N elements
  - How many reallocations will happen?
  - How many moves/copies will happen?
  - Is it amortized constant complexity?
    - That means the ratio between number of moves/copies and N is constant.



# push\_back complexity analysis


$$\text{numRealloc} = \lceil \log_k N \rceil$$

$$N = 10, k = 2 \Rightarrow \text{numRealloc} = 4$$

$$\text{numMovCp} = k^0 + k^1 + k^2 + \dots + k^{\text{numRealloc}-1}$$

$$= \frac{k^{\text{numRealloc}} - 1}{k - 1} = \frac{k^{\lceil \log_k N \rceil} - 1}{k - 1}$$

[https://en.wikipedia.org/wiki/Geometric\\_series](https://en.wikipedia.org/wiki/Geometric_series)




# push\_back complexity analysis

$$\log_k N \leq \lceil \log_k N \rceil < \log_k N + 1$$

$$\Rightarrow N \leq k^{\lceil \log_k N \rceil} < Nk$$

$$\Rightarrow \frac{N-1}{k-1} \leq \text{numMovCp} < \frac{Nk-1}{k-1}$$

$$\Rightarrow \frac{1 - \frac{1}{N}}{k-1} \approx C_1 \leq \frac{\text{numMovCp}}{N} < \frac{k - \frac{1}{N}}{k-1} \approx C_2$$



# Reallocation


- How can we move/copy old elements?
  - Move
    - Efficient.
    - Not everything can be moved.
    - Move is not always noexcept.
  - Copy
    - Inefficient (imagine we're copying 1'000 std::string).
    - Strong exception guarantee.
- Move when noexcept, copy otherwise.
  - If type is uncopyable, and move is not noexcept, still use move.
  - Dangerous.

# emplace\_back

- Complexity: Amortized constant.
- Exception: Strong.
- Implementation
  - Use perfect forwarding.
  - `emplace_back` constructs object inplace.
- Can we `emplace_back` a self element?  

```
vector<int> v{1, 2, 3};  
v.emplace_back(v[0]);
```

  - Yes.



# push\_back

- Complexity: Amortized constant.
- Exception: Strong.
- Implementation
  - Copy or move the input object into vector.
  - Can be costly.
- Can we push\_back a self element?
  - Yes.

# reserve

- Complexity: Usually  $O(N)$
- Exception: Strong
- Implementation
  - Reallocate new memory.
  - Move/copy old elements.
- Used to enlarge “unused” space.

# What's wrong with this?

```
template <class T> struct my_ptr {  
    my_ptr(T* p) : m_p{p} {}  
    ~my_ptr() { delete m_p; }  
};  
vector<my_ptr<int>> v;  
v.push_back(new int{0});  
v.emplace_back(new int{0});
```

- Don't use new, unless we have a reason.



# What's wrong with this?

```
vector<int> v;  
for (auto i = 0; i < 1000; ++i)  
    v.emplace_back(i);
```

- We should do a reserve first.

# What's wrong with this?

```
vector<int> v{1, 2, 3};  
auto x = v.at(0);
```

- `vector::at` is useless, use `vector::operator[]` instead.

# Recommendation

---

- Design your type so it has noexcept move.
- You should use `std::vector` by default.
- "Use vectors whenever you can. If you cannot use vectors, redesign your solution so that you can use vectors." - Alexander Stepanov

Thanks for listening.

Question?