

The STL design: an overview

HCM City C++ users meetup

Germán Diago Gómez

May 22nd, 2016

Goals of this talk

- introduce STL fundamental ideas

Goals of this talk

- introduce STL fundamental ideas
- understand the design choices that drove its current design

Goals of this talk

- introduce STL fundamental ideas
- understand the design choices that drove its current design
- light introduction to each of its fundamental parts

Goals of this talk

- introduce STL fundamental ideas
- understand the design choices that drove its current design
- light introduction to each of its fundamental parts

Goals of this talk

- introduce STL fundamental ideas
- understand the design choices that drove its current design
- light introduction to each of its fundamental parts

Non-goals

- not a detailed, fine-grained overview

Non-goals

- not a detailed, fine-grained overview
- not lots of examples

Non-goals

- not a detailed, fine-grained overview
- not lots of examples
 - due to time constraints

STL

- STL stands for Standard Template Library

STL

- STL stands for Standard Template Library
- it is a subset of the C++ Standard Library

STL

- STL stands for Standard Template Library
- it is a subset of the C++ Standard Library
- its main author is Alexander Stepanov

STL

- STL stands for Standard Template Library
- it is a subset of the C++ Standard Library
- its main author is Alexander Stepanov
- got accepted into ISO C++98 standard library

Design

- general

Design

- general
- allows a high degree of customization through composition

Design

- general
- allows a high degree of customization through composition
- extensible: you can add new components to the framework

Design

- general
- allows a high degree of customization through composition
- extensible: you can add new components to the framework
- reusable **without loss of performance** compared to hand-written code

Design

- general
- allows a high degree of customization through composition
- extensible: you can add new components to the framework
- reusable **without loss of performance** compared to hand-written code
- intensive use of generic programming

STL main components

- containers

STL main components

- containers
- iterators

STL main components

- containers
- iterators
- algorithms

STL main components

- containers
- iterators
- algorithms
- utilities (not the focus of this talk)

Why generic programming

- allows compile-time polymorphism combined with overloading

Why generic programming

- allows compile-time polymorphism combined with overloading
- allows customized code generation at compile-time

Why generic programming

- allows compile-time polymorphism combined with overloading
- allows customized code generation at compile-time
- allows choose best optimization based on overloading, which is a compile-time mechanism

Why generic programming

- allows compile-time polymorphism combined with overloading
- allows customized code generation at compile-time
- allows choose best optimization based on overloading, which is a compile-time mechanism
- allows high degree of customization, for both algorithms and data structures **without loss of performance**

Why generic programming (2)

- types can retroactively model a concept

Why generic programming (2)

- types can retroactively model a concept
 - example: array pointers are iterators

Why generic programming (2)

- types can retroactively model a concept
 - example: array pointers are iterators
- OO alternative: uses run-time polymorphism and dispatch

Why generic programming (2)

- types can retroactively model a concept
 - example: array pointers are iterators
- OO alternative: uses run-time polymorphism and dispatch
 - performance penalty

Concepts briefly explained

- Concepts are the foundation of generic programming

Concepts briefly explained

- Concepts are the foundation of generic programming
- a Concept is a set of requirements that a type must meet

Concepts briefly explained

- Concepts are the foundation of generic programming
- a Concept is a set of requirements that a type must meet
 - the **syntactic** requirements that a type must meet

Concepts briefly explained

- Concepts are the foundation of generic programming
- a Concept is a set of requirements that a type must meet
 - the **syntactic** requirements that a type must meet
 - the **semantic** requirements that these syntactic requirements must meet

Concepts briefly explained

- Concepts are the foundation of generic programming
- a Concept is a set of requirements that a type must meet
 - the **syntactic** requirements that a type must meet
 - the **semantic** requirements that these syntactic requirements must meet
 - optionally, **associated types** for that concept

Concepts briefly explained

- Concepts are the foundation of generic programming
- a Concept is a set of requirements that a type must meet
 - the **syntactic** requirements that a type must meet
 - the **semantic** requirements that these syntactic requirements must meet
 - optionally, **associated types** for that concept
 - algorithms: include the big O cost associated to them

Concepts briefly explained

- Concepts are the foundation of generic programming
- a Concept is a set of requirements that a type must meet
 - the **syntactic** requirements that a type must meet
 - the **semantic** requirements that these syntactic requirements must meet
 - optionally, **associated types** for that concept
 - algorithms: include the big O cost associated to them
- In C++ these are the requirements that your template parameters need to fulfill. They are stated in the documentation, since there is no language support (yet) for Concepts

Concepts briefly explained (2)

- similar to the way a virtual function sets requirements on derived types

Concepts briefly explained (2)

- similar to the way a virtual function sets requirements on derived types
 - but at compile-time

Concepts briefly explained (2)

- similar to the way a virtual function sets requirements on derived types
 - but at compile-time
- when a type fulfills the requirements of a concept we say it **models** that concept

Concepts briefly explained (2)

- similar to the way a virtual function sets requirements on derived types
 - but at compile-time
- when a type fulfills the requirements of a concept we say it **models** that concept
 - similar to when we say a derived class **implements** requirements from a base class

Concepts briefly explained (2)

- similar to the way a virtual function sets requirements on derived types
 - but at compile-time
- when a type fulfills the requirements of a concept we say it **models** that concept
 - similar to when we say a derived class **implements** requirements from a base class
- types do not need to inherit from any class to model a Concept → more loosely coupled

Concepts example: LessThanComparable

- syntactic requirements

Concepts example: LessThanComparable

- syntactic requirements
 - your type must work with `operator<` and return a type convertible to `bool`

Concepts example: LessThanComparable

- syntactic requirements
 - your type must work with `operator<` and return a type convertible to `bool`
- semantic requirements for `LessThanComparable` `operator<`

Concepts example: LessThanComparable

- syntactic requirements
 - your type must work with `operator<` and return a type convertible to `bool`
- semantic requirements for `LessThanComparable` `operator<`
 - $\forall a, !(a < a)$

Concepts example: LessThanComparable

- syntactic requirements
 - your type must work with `operator<` and return a type convertible to `bool`
- semantic requirements for `LessThanComparable` `operator<`
 - $\forall a, !(a < a)$
 - $a < b \iff !(b < a)$

Concepts example: LessThanComparable

- syntactic requirements
 - your type must work with `operator<` and return a type convertible to `bool`
- semantic requirements for `LessThanComparable` `operator<`
 - $\forall a, \neg(a < a)$
 - $a < b \iff \neg(b < a)$
 - $a < b \wedge b < c \implies a < c$

Concepts example: LessThanComparable

- syntactic requirements
 - your type must work with `operator<` and return a type convertible to `bool`
- semantic requirements for `LessThanComparable` `operator<`
 - $\forall a, !(a < a)$
 - $a < b \iff !(b < a)$
 - $a < b \wedge b < c \implies a < c$
- associated types (optionally)

More Concepts examples

- BinaryPredicate
- DefaultConstructible
- Many others

Introduction

Containers allow you to store collections of data in different data structures that best suite your needs.

STL containers classification

- sequence containers

STL containers classification

- sequence containers
 - `std::array<T>`, `std::vector<T>`, `std::list<T>`...

STL containers classification

- sequence containers
 - `std::array<T>`, `std::vector<T>`, `std::list<T>`...
- associative containers

STL containers classification

- sequence containers
 - `std::array<T>`, `std::vector<T>`, `std::list<T>`...
- associative containers
 - ordered \rightarrow `std::set`...

STL containers classification

- sequence containers
 - `std::array<T>`, `std::vector<T>`, `std::list<T>`...
- associative containers
 - ordered \rightarrow `std::set`...
 - unordered \rightarrow `std::unordered_map`...

STL containers classification

- sequence containers
 - `std::array<T>`, `std::vector<T>`, `std::list<T>`...
- associative containers
 - ordered \rightarrow `std::set`...
 - unordered \rightarrow `std::unordered_map`...
- container adaptors

STL containers classification

- sequence containers
 - `std::array<T>`, `std::vector<T>`, `std::list<T>`...
- associative containers
 - ordered \rightarrow `std::set`...
 - unordered \rightarrow `std::unordered_map`...
- container adaptors
 - `std::priority_queue`, `std::stack`...

STL containers traversal

- depends on the kind iterators provided (more later)

STL containers traversal

- depends on the kind iterators provided (more later)
- random access, bidirectional, forward containers, etc.

Introduction

- Iterators are used to traverse ranges

Introduction

- Iterators are used to traverse ranges
- Iterators are an abstraction for positions

Introduction

- Iterators are used to traverse ranges
- Iterators are an abstraction for positions
 - they are modeled after C pointers in C++'s concrete case

Introduction

- Iterators are used to traverse ranges
- Iterators are an abstraction for positions
 - they are modeled after C pointers in C++'s concrete case
- you can do a traversal with **two** iterators

Introduction

- Iterators are used to traverse ranges
- Iterators are an abstraction for positions
 - they are modeled after C pointers in C++'s concrete case
- you can do a traversal with **two** iterators
- they are **not** the same as an iterator in other languages

Introduction

- Iterators are used to traverse ranges
- Iterators are an abstraction for positions
 - they are modeled after C pointers in C++'s concrete case
- you can do a traversal with **two** iterators
- they are **not** the same as an iterator in other languages
 - Java/C# iterator similar to a **pair of iterators** in C++

Introduction

- Iterators are used to traverse ranges
- Iterators are an abstraction for positions
 - they are modeled after C pointers in C++'s concrete case
- you can do a traversal with **two** iterators
- they are **not** the same as an iterator in other languages
 - Java/C# iterator similar to a **pair of iterators** in C++
- a pair of iterators is usually called a **range** in C++

Iterator categories

- several kinds of iterators

Iterator categories

- several kinds of iterators
- iterators form a hierarchy of Concepts

Iterator categories

- several kinds of iterators
- iterators form a hierarchy of Concepts
 - not implemented as inheritance in the library

Iterator categories

- several kinds of iterators
- iterators form a hierarchy of Concepts
 - not implemented as inheritance in the library
 - they are Concepts

Iterator categories

- several kinds of iterators
- iterators form a hierarchy of Concepts
 - not implemented as inheritance in the library
 - they are Concepts
- `RandomAccessIterator` → `BidirectionalIterator` → `ForwardIterator` → `InputIterator`

Iterator categories

- several kinds of iterators
- iterators form a hierarchy of Concepts
 - not implemented as inheritance in the library
 - they are Concepts
- `RandomAccessIterator` \rightarrow `BidirectionalIterator` \rightarrow `ForwardIterator` \rightarrow `InputIterator`
- `OutputIterator` is any iterator that can write to the element it points to

Iterator categories (2)

- every iterator category supports a set of operations

Iterator categories (2)

- every iterator category supports a set of operations
- InputIterator

Iterator categories (2)

- every iterator category supports a set of operations
- InputIterator
 - element access

Iterator categories (2)

- every iterator category supports a set of operations
- InputIterator
 - element access
 - increment iterator position by one

Iterator categories (2)

- every iterator category supports a set of operations
- InputIterator
 - element access
 - increment iterator position by one
 - single traversal

Iterator categories (2)

- every iterator category supports a set of operations
- InputIterator
 - element access
 - increment iterator position by one
 - single traversal
- ForwardIterator

Iterator categories (2)

- every iterator category supports a set of operations
- InputIterator
 - element access
 - increment iterator position by one
 - single traversal
- ForwardIterator
 - InputIterator + support for multiple traversals

Iterator categories (2)

- every iterator category supports a set of operations
- InputIterator
 - element access
 - increment iterator position by one
 - single traversal
- ForwardIterator
 - InputIterator + support for multiple traversals
- BidirectionalIterator

Iterator categories (2)

- every iterator category supports a set of operations
- InputIterator
 - element access
 - increment iterator position by one
 - single traversal
- ForwardIterator
 - InputIterator + support for multiple traversals
- BidirectionalIterator
 - ForwardIterator + supports decrementing position by one

Iterator categories (2)

- every iterator category supports a set of operations
- InputIterator
 - element access
 - increment iterator position by one
 - single traversal
- ForwardIterator
 - InputIterator + support for multiple traversals
- BidirectionalIterator
 - ForwardIterator + supports decrementing position by one
- RandomAccessIterator

Iterator categories (2)

- every iterator category supports a set of operations
- InputIterator
 - element access
 - increment iterator position by one
 - single traversal
- ForwardIterator
 - InputIterator + support for multiple traversals
- BidirectionalIterator
 - ForwardIterator + supports decrementing position by one
- RandomAccessIterator
 - BidirectionalIterator + supports jumping to arbitrary position in $O(1)$

Iterator categories (2)

- every iterator category supports a set of operations
- InputIterator
 - element access
 - increment iterator position by one
 - single traversal
- ForwardIterator
 - InputIterator + support for multiple traversals
- BidirectionalIterator
 - ForwardIterator + supports decrementing position by one
- RandomAccessIterator
 - BidirectionalIterator + supports jumping to arbitrary position in $O(1)$
 - think of a pointer to C array position

Why are iterators important

- abstract traversal from the underlying data structure

Why are iterators important

- abstract traversal from the underlying data structure
- decouple algorithms implementation from data structure (later more)

Why are iterators important

- abstract traversal from the underlying data structure
- decouple algorithms implementation from data structure (later more)
- can implement algorithms based on its iterator category (later more)

Why are iterators important

- abstract traversal from the underlying data structure
- decouple algorithms implementation from data structure (later more)
- can implement algorithms based on its iterator category (later more)
 - containers just need to provide a begin/end sequence pair

Why are iterators important

- abstract traversal from the underlying data structure
- decouple algorithms implementation from data structure (later more)
- can implement algorithms based on its iterator category (later more)
 - containers just need to provide a begin/end sequence pair
- let you reduce implementation effort on algorithms

Introduction

- algorithms let you perform operations on your data

Introduction

- algorithms let you perform operations on your data
- STL algorithms design choices are a bit suprising at first

Algorithm design choices

- it was decided to implement algorithms as free functions, not as member functions

Algorithm design choices

- it was decided to implement algorithms as free functions, not as member functions
- it was decided that it would take `Iterator` s as input, instead of `Container` s

Algorithm design choices

- it was decided to implement algorithms as free functions, not as member functions
- it was decided that it would take `Iterator` s as input, instead of `Container` s
- consequence: algorithms are decoupled from containers

STL algorithms/iterators conventions

- algorithm specifications use half-open ranges for algorithms

STL algorithms/iterators conventions

- algorithm specifications use half-open ranges for algorithms
 - $[a, b)$

STL algorithms/iterators conventions

- algorithm specifications use half-open ranges for algorithms
 - $[a, b)$
 - a is included in the range. b is **not** included in the range

STL algorithms/iterators conventions

- algorithm specifications use half-open ranges for algorithms
 - $[a, b)$
 - a is included in the range. b is **not** included in the range
- containers provide `begin`, `end` and its variants

STL algorithms/iterators conventions

- algorithm specifications use half-open ranges for algorithms
 - $[a, b)$
 - a is included in the range. b is **not** included in the range
- containers provide `begin`, `end` and its variants
 - `begin` points to the first element of the container

STL algorithms/iterators conventions

- algorithm specifications use half-open ranges for algorithms
 - $[a, b)$
 - a is included in the range. b is **not** included in the range
- containers provide `begin`, `end` and its variants
 - `begin` points to the first element of the container
 - `end` points **one past the last element** of the container traversal

STL algorithms/iterators conventions

- algorithm specifications use half-open ranges for algorithms
 - $[a, b)$
 - a is included in the range. b is **not** included in the range
- containers provide `begin`, `end` and its variants
 - `begin` points to the first element of the container
 - `end` points **one past the last element** of the container traversal
- thus, passing `begin` and `end` to an algorithm will traverse the whole sequence

STL algorithms/iterators conventions

- algorithm specifications use half-open ranges for algorithms
 - $[a, b)$
 - a is included in the range. b is **not** included in the range
- containers provide `begin`, `end` and its variants
 - `begin` points to the first element of the container
 - `end` points **one past the last element** of the container traversal
- thus, passing `begin` and `end` to an algorithm will traverse the whole sequence
- some algorithms take three iterators: a range + the beginning of a second range where to put the output

STL algorithms/iterators conventions

- algorithm specifications use half-open ranges for algorithms
 - $[a, b)$
 - a is included in the range. b is **not** included in the range
- containers provide `begin`, `end` and its variants
 - `begin` points to the first element of the container
 - `end` points **one past the last element** of the container traversal
- thus, passing `begin` and `end` to an algorithm will traverse the whole sequence
- some algorithms take three iterators: a range + the beginning of a second range where to put the output
 - C++14: these algorithms, called "three-legged" algorithms, were added overloads to pass two ranges. Example: `std::equal`

STL algorithms/iterators conventions (2)

- no memory management in algorithms. Beware

STL algorithms/iterators conventions (2)

- no memory management in algorithms. Beware
 - you must have space in the target output range

STL algorithms/iterators conventions (2)

- no memory management in algorithms. Beware
 - you must have space in the target output range
 - example: `std::copy` must have enough space for the output

STL algorithms/iterators conventions (2)

- no memory management in algorithms. Beware
 - you must have space in the target output range
 - example: `std::copy` must have enough space for the output
 - `std::remove` algorithms do not resize your containers

STL algorithms/iterators conventions (2)

- no memory management in algorithms. Beware
 - you must have space in the target output range
 - example: `std::copy` must have enough space for the output
 - `std::remove` algorithms do not resize your containers
 - they just put the rubbish to remove at the end of your container. You must call `Container::erase`.

STL algorithms/iterators conventions (2)

- no memory management in algorithms. Beware
 - you must have space in the target output range
 - example: `std::copy` must have enough space for the output
 - `std::remove` algorithms do not resize your containers
 - they just put the rubbish to remove at the end of your container. You must call `Container::erase`.
- you can use iterator adaptors that make room for elements if you do not know how much space you will need. example: `back_insert_iterator`.

Extensibility and adaptation

- you can provide an algorithm based on iterator category: it will work on all existing containers that model at least that iterator category

Extensibility and adaptation

- you can provide an algorithm based on iterator category: it will work on all existing containers that model at least that iterator category
- if you provide your iterators for you container, it just works

Extensibility and adaptation

- you can provide an algorithm based on iterator category: it will work on all existing containers that model at least that iterator category
- if you provide your iterators for you container, it just works
- basic arrays also work with algorithms

Implementation effort reduction

- container-based implementation: N algorithms for M data structures

Implementation effort reduction

- container-based implementation: N algorithms for M data structures
 - $O(N \cdot M)$ implementations

Implementation effort reduction

- container-based implementation: N algorithms for M data structures
 - $O(N \cdot M)$ implementations
- iterator-based implementation: N algorithms independent of the data structure

Implementation effort reduction (2)

- STL \approx 60 algorithms and 13 data structures (without counting adapters and native arrays)

Implementation effort reduction (2)

- STL \approx 60 algorithms and 13 data structures (without counting adapters and native arrays)
 - container-based algorithms: $60 * 13 = 780$ implementations

Implementation effort reduction (2)

- STL \approx 60 algorithms and 13 data structures (without counting adapters and native arrays)
 - container-based algorithms: $60 * 13 = 780$ implementations
 - iterator-based algorithms: 60 implementations, plus maybe some optimization overloads in cases where algorithms can be optimized for a specific iterator category

Implementation effort reduction (2)

- STL \approx 60 algorithms and 13 data structures (without counting adapters and native arrays)
 - container-based algorithms: $60 * 13 = 780$ implementations
 - iterator-based algorithms: 60 implementations, plus maybe some optimization overloads in cases where algorithms can be optimized for a specific iterator category
 - remember that algorithms must be as good as hand-written algorithms for any use case

STL algorithm example

Utilities

Utilities include other useful components such as `std::pair`, `std::function`, `type_traits`, functional objects, `std::tuple`, `std::unique_ptr`, `std::shared_ptr`, etc.

Typical algorithm usage

- 1 Choose an algorithm.

Typical algorithm usage

- 1 Choose an algorithm.
- 2 Provide the iterators from your container(s). for the algorithm through `begin()` and `end()` functions

Typical algorithm usage

- 1 Choose an algorithm.
- 2 Provide the iterators from your container(s). for the algorithm through `begin()` and `end()` functions
- 3 Customize the behaviour of your algorithm through predicates/function objects.

Example exercise

Read all the data for the spanish football league. We store the data in a `league` class which is a collection of `football_match_result` records. The file is sorted by round. Every 10 matches it is a new round. Blank lines are accepted.

Example exercise

Read all the data for the spanish football league. We store the data in a `league` class which is a collection of `football_match_result` records. The file is sorted by round. Every 10 matches it is a new round. Blank lines are accepted.

File `spanish_football_league2015-2016.txt`

```
Real Madrid,2-1,Malaga  
Rayo,0-2,Sevilla  
...
```

Example exercise

Read all the data for the spanish football league. We store the data in a `league` class which is a collection of `football_match_result` records. The file is sorted by round. Every 10 matches it is a new round. Blank lines are accepted.

File `spanish_football_league2015-2016.txt`

```
Real Madrid,2-1,Malaga  
Rayo,0-2,Sevilla  
...
```

We will:

- 1 Get all matches for a given team

Example exercise

Read all the data for the spanish football league. We store the data in a `league` class which is a collection of `football_match_result` records. The file is sorted by round. Every 10 matches it is a new round. Blank lines are accepted.

File `spanish_football_league2015-2016.txt`

```
Real Madrid,2-1,Malaga  
Rayo,0-2,Sevilla  
...
```

We will:

- 1 Get all matches for a given team
- 2 Calculate the classification table

Example exercise

Read all the data for the spanish fooball league. We store the data in a `league` class which is a collection of `football_match_result` records. The file is sorted by round. Every 10 matches it is a new round. Blank lines are accepted.

File `spanish_football_league2015-2016.txt`

```
Real Madrid,2-1,Malaga  
Rayo,0-2,Sevilla  
...
```

We will:

- 1 Get all matches for a given team
- 2 Calculate the classification table
- 3 Calculate the matches a given team won as a visitor

Current status of future standardization

- added parallel algorithms to C++17
 - essentially, you do `std::sort(parallel::par, begin(v), end(v))` and enjoy performance gains automatically
- ranges-based STL, by Eric Niebler → Ranges v3 experimental implementation
 - you can do `std::sort(v)` and much more
 - lazy evaluation and composition
 - `v | view::transform([](auto const & elem) {...}) | view::remove_if([]...)`
 - library includes, besides ranges, view s and action s, but these are not proposed for standardization yet
- there is a TS for Concepts support in the language: Concepts TS latest draft
 - implementation available in GCC 6 with `-fconcepts` flag

Thank you